

MCMC software

Stat 340: Bayesian Statistics



**RUN
MCMC**

Implementing MCMC in JAGS

(Problem topic 9)

Example

- Researchers are interested in modeling the survival times, measured in weeks, of patients who were diagnosed with leukemia.
- The patients were classified according to two characteristics of white blood cells.
- The sample consists of $n = 17$ times in weeks from diagnosis to death

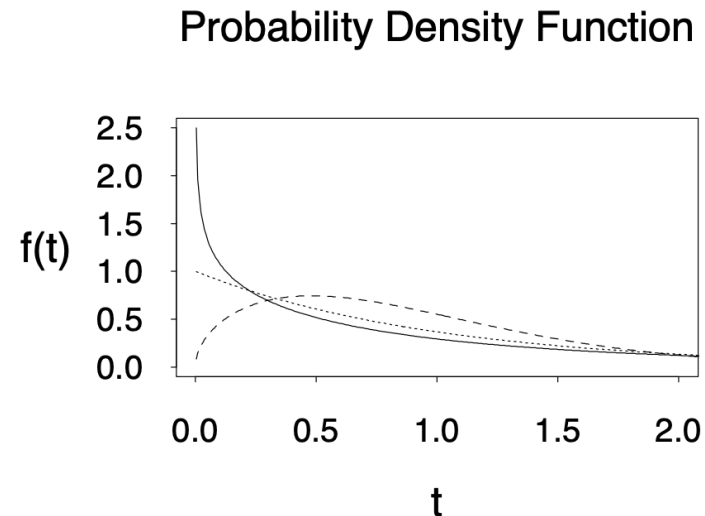
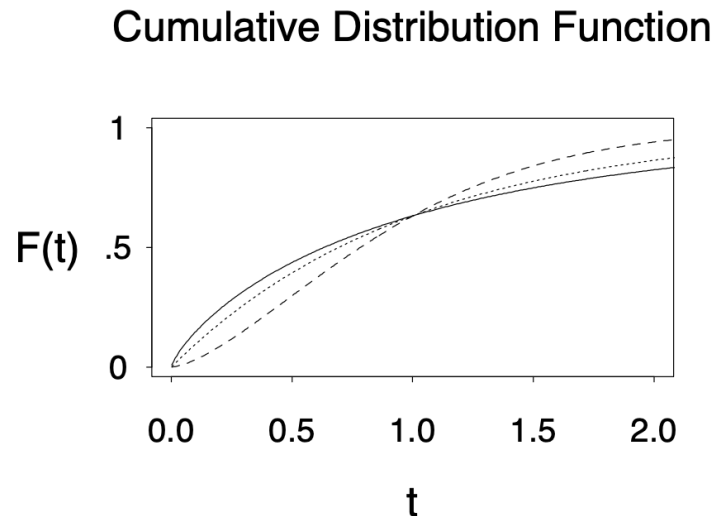
```
survival_times <- c(65, 156, 100, 134, 16, 108, 121, 4,  
                   39, 143, 56, 26, 22, 1, 1, 5, 65)
```

Weibull distribution

Weibull distribution is often used as a model for survival times in biomedical, demographic, and engineering analyses

PDF: $f(y|v, \lambda) = \lambda \alpha y^{\alpha-1} e^{-\lambda y^\alpha}, \quad y, \alpha, \lambda > 0$

CDF: $F(y|v, \lambda) = 1 - e^{-\lambda y^\alpha}$



Proposed model

Let Y_i = survival time in weeks


$$Y_i | \alpha, \lambda \stackrel{\text{iid}}{\sim} \text{Weibull}(\alpha, \lambda)$$

$$\alpha \sim \text{Gamma}(1, 1)$$

$$\lambda \sim \text{Gamma}(1.53, 26.3)$$

$$\alpha \propto \lambda$$

Prior predictive



Your turn

Weibull PDF: $f(y|v, \lambda) = \lambda \alpha y^{\alpha-1} e^{-\lambda y^\alpha}$

Gamma PDF: $f(x|a, b) = \frac{b^a}{\Gamma(a)} x^{a-1} e^{-bx}$

1. Derive the joint posterior distribution for α and λ
2. Derive the two conditional posterior distributions for this model
3. Are either distributions that you recognize?
4. If you get done, outline the steps of a Gibbs sampler for this posterior

05:00

Why Just Another Gibbs Sampler (JAGS)?

- You can fit virtually any model
- You can call JAGS from R, allowing for plotting and data manipulation in R
- It runs on all platforms: LINUX, Mac, Windows
- There is a lot of help online
- R has many built in packages for convergence diagnostics

How does JAGS work?

- You specify the model by declaring the likelihood and priors
- JAGS then sets up the MCMC sampler
 - e.g., works out the full conditional distributions for all parameters
- It returns MCMC samples in a matrix or array
- It also automatically produces posterior summaries like means, credible sets, and convergence diagnostics
- User's manual:
http://people.stat.sc.edu/hansont/stat740/jags_user_manual.pdf

Steps for any analysis using JAGS

1. Specify the model as a string (or script)
2. Define/load the data
3. Define initial values
4. Draw posterior samples using `run.jags()`
5. Inspect/summarize the results

0. Load `runjags`

```
library(runjags)
```

1. Specify the model as a string

```
model_string <- "  
model{  
  # Specify the likelihood  
  for(i in 1:n) {  
    y[i] ~ dweib(alpha, lambda)  
  }  
  
  # Specify the priors  
  alpha ~ dgamma(1, 1)  
  lambda ~ dgamma(1.53, 26.3)  
}  
"
```

JAGS syntax

JAGS doesn't always use the same syntax as R.

But, uses the same **d***, **q***, and **p*** prefixes

Distribution	Distribution	Quantile	
Bernoulli	dbern	pbern	qbern
Beta	dbeta	pbeta	qbeta
Binomial	dbin	pbin	qbin
Chi-squared	dchisqr	pchisqr	qchisqr
Double exponential	ddexp	pdexp	qdexp
Exponential	dexp	pexp	qexp
F	df	pf	qf
Gamma	dgamma	pgamma	qgamma
Generalized gamma	dgen.gamma	pgen.gamma	qgen.gamma
Noncentral hypergeometric	dhyper	phyper	qhyper
Logistic	dlogis	plogis	qlogis
Log-normal	dlnorm	plnorm	qlnorm
Negative binomial	dnegbin	pnegbin	qnegbin
Noncentral Chi-squared	dnchisqr	pnchisqr	qnchisqr
Normal	dnorm	pnorm	qnorm
Pareto	dpar	ppar	qpar
Poisson	dpois	ppois	qpois
Student t	dt	pt	qt
Weibull	dweib	pweib	qweib

Image source: JAGS manual

2. Define/load the data

Loading the survival data

```
survival_times <- c(65, 156, 100, 134, 16, 108, 121, 4,  
                   39, 143, 56, 26, 22, 1, 1, 5, 65)
```

Store your data in a list

```
survival_data <- list(  
  y = survival_times,  
  n = length(survival_times)  
)
```

3. Define initial values

- If you don't specify initial values for your parameters, then JAGS will
- If not specified, then JAGS will use the mean or mode of the prior distribution
- More on this next time...

4. Draw posterior samples using `run.jags()`

- `model` string specifying the model

```
posterior <- run.jags(  
  model = model_string,  
  n.chains = 1,  
  data = survival_data,  
  monitor = c("alpha", "lambda"),  
  adapt = 1000,  
  burnin = 5000,  
  sample = 5000  
)
```

4. Draw posterior samples using `run.jags()`

- `model` string specifying the model
- `n.chains` the number of chains to run

```
posterior <- run.jags(  
  model = model_string,  
  n.chains = 1,  
  data = survival_data,  
  monitor = c("alpha", "lambda"),  
  adapt = 1000,  
  burnin = 5000,  
  sample = 5000  
)
```


4. Draw posterior samples using `run.jags()`

- `model` string specifying the model
- `n.chains` the number of chains to run
- `data` a named list or data frame include the data and prior parameter values

```
posterior <- run.jags(  
  model = model_string,  
  n.chains = 1,  
  data = survival_data,  
  monitor = c("alpha", "lambda"),  
  adapt = 1000,  
  burnin = 5000,  
  sample = 5000  
)
```

4. Draw posterior samples using `run.jags()`

- `model` string specifying the model
- `n.chains` the number of chains to run
- `data` a named list or data frame include the data and prior parameter values
- `monitor` character vector of the names of variables to monitor

```
posterior <- run.jags(  
  model = model_string,  
  n.chains = 1,  
  data = survival_data,  
  monitor = c("alpha", "lambda"),  
  adapt = 1000,  
  burnin = 5000,  
  sample = 5000  
)
```

4. Draw posterior samples using `run.jags()`

- `model` string specifying the model
- `n.chains` the number of chains to run
- `data` a named list or data frame include the data and prior parameter values
- `monitor` character vector of the names of variables to monitor
- `adapt` number of samples drawn during initial sampling/adaptation phase

```
posterior <- run.jags(  
  model = model_string,  
  n.chains = 1,  
  data = survival_data,  
  monitor = c("alpha", "lambda"  
  adapt = 1000,  
  burnin = 5000,  
  sample = 5000  
)
```

4. Draw posterior samples using `run.jags()`

- `model` string specifying the model
- `n.chains` the number of chains to run
- `data` a named list or data frame include the data and prior parameter values
- `monitor` character vector of the names of variables to monitor
- `adapt` number of samples drawn during initial sampling/adaptation phase
- `burnin` number of burn-in iterations, NOT including the adaptive iterations

```
posterior <- run.jags(  
  model = model_string,  
  n.chains = 1,  
  data = survival_data,  
  monitor = c("alpha", "lambda"),  
  adapt = 1000,  
  burnin = 5000,  
  sample = 5000  
)
```

4. Draw posterior samples using `run.jags()`

- `model` string specifying the model
- `n.chains` the number of chains to run
- `data` a named list or data frame include the data and prior parameter values
- `monitor` character vector of the names of variables to monitor
- `adapt` number of samples drawn during initial sampling/adaptation phase
- `burnin` number of burn-in iterations, NOT including the adaptive iterations
- `sample` the total number of (additional) samples to draw

```
posterior <- run.jags(  
  model = model_string,  
  n.chains = 1,  
  data = survival_data,  
  monitor = c("alpha", "lambda"),  
  adapt = 1000,  
  burnin = 5000,  
  sample = 5000  
)
```









In an .Rmd file, add `silent.jags = TRUE` to avoid this issue...

```
## Compiling rjags model...  
## Calling the simulation using the rjags method...  
## Adapting the model for 1000 iterations...  
## Burning in the model for 5000 iterations...  
## Running the model for 5000 iterations...  
## Simulation complete  
## Calculating summary statistics...
```

```
## Warning: Convergence cannot be assessed with only 1 chain
```

```
## Finished running the simulation
```

5. Inspect the results



runjags objects are big lists

```
names(posterior)
```

```
## [1] "mcmc" "deviance.table" "deviance.sum"
## [4] "pd" "end.state" "samplers"
## [7] "burnin" "sample" "thin"
## [10] "model" "data" "monitor"
## [13] "noread.monitor" "modules" "factories"
## [16] "response" "residual" "fitted"
## [19] "method" "method.options" "timetaken"
## [22] "runjags.version" "summaries" "summary"
## [25] "HPD" "hpd" "mcse"
## [28] "psrf" "autocorr" "crosscorr"
## [31] "truestochastic" "semistochastic" "nonstochastic"
## [34] "discrete" "trace" "density"
## [37] "histogram" "ecdfplot" "key"
## [40] "acplot" "ccplot" "summary.available"
## [43] "summary.pars" "dic"
```


`mcmc` element is a list of the actual draws

```
head(posterior$mcmc)
```

```
## [[1]]  
## Markov Chain Monte Carlo (MCMC) output:  
## Start = 6001  
## End = 6007  
## Thinning interval = 1  
##           alpha      lambda  
## 6001 0.8868662 0.01977346  
## 6002 0.9252918 0.01395950  
## 6003 1.0516895 0.01721196  
## 6004 0.9212670 0.02472420  
## 6005 0.9250907 0.02964864  
## 6006 0.7915493 0.05146866  
## 6007 0.7406706 0.05755608  
##  
## attr(,"class")  
## [1] "mcmc.list"
```

5. Summarize the results

```
summary(posterior$mcmc[[1]])  
##  
## Iterations = 6001:11000  
## Thinning interval = 1  
## Number of chains = 1  
## Sample size per chain = 5000  
##  
## 1. Empirical mean and standard deviation for each variable,  
##    plus standard error of the mean:  
##  
##           Mean      SD Naive SE Time-series SE  
## alpha  0.82868 0.13494 0.0019084      0.007950  
## lambda 0.04096 0.02484 0.0003512      0.001279  
##  
## 2. Quantiles for each variable:  
##  
##           2.5%      25%      50%      75%      97.5%  
## alpha  0.58207 0.73092 0.82469 0.9219 1.1017  
## lambda 0.00974 0.02251 0.03521 0.0537 0.1027
```

bayesplot

- provides a variety of `ggplot2`-based plotting functions for use after fitting Bayesian models
- Visualizations of MCMC simulations and diagnostics *from any algorithm* (`mcmc_*`)
- Graphical prior and posterior predictive checks (`ppc_*`)

Convert to an `mcmc` object

`bayesplot` requires objects to be of class `mcmc`


You'll need to convert your output from `run.jags()`

```
class(posterior)
## [1] "runjags"
```

```
post_mcmc <- as.mcmc(posterior)
class(post_mcmc)
## [1] "mcmc"
```

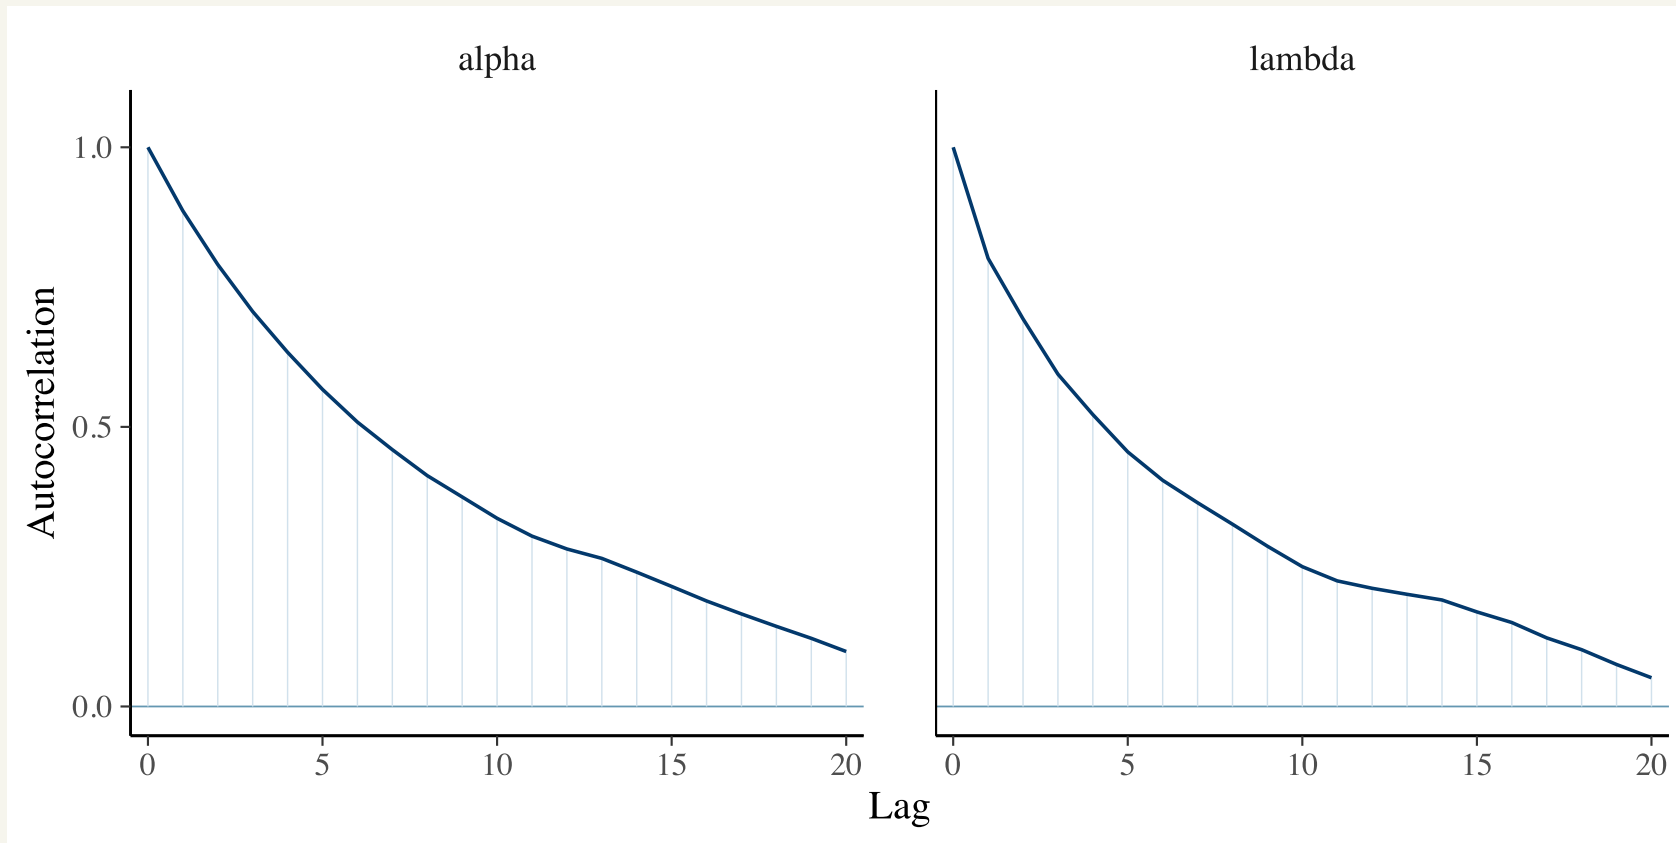
Trace plot

```
mcmc_trace(post_mcmc)
```



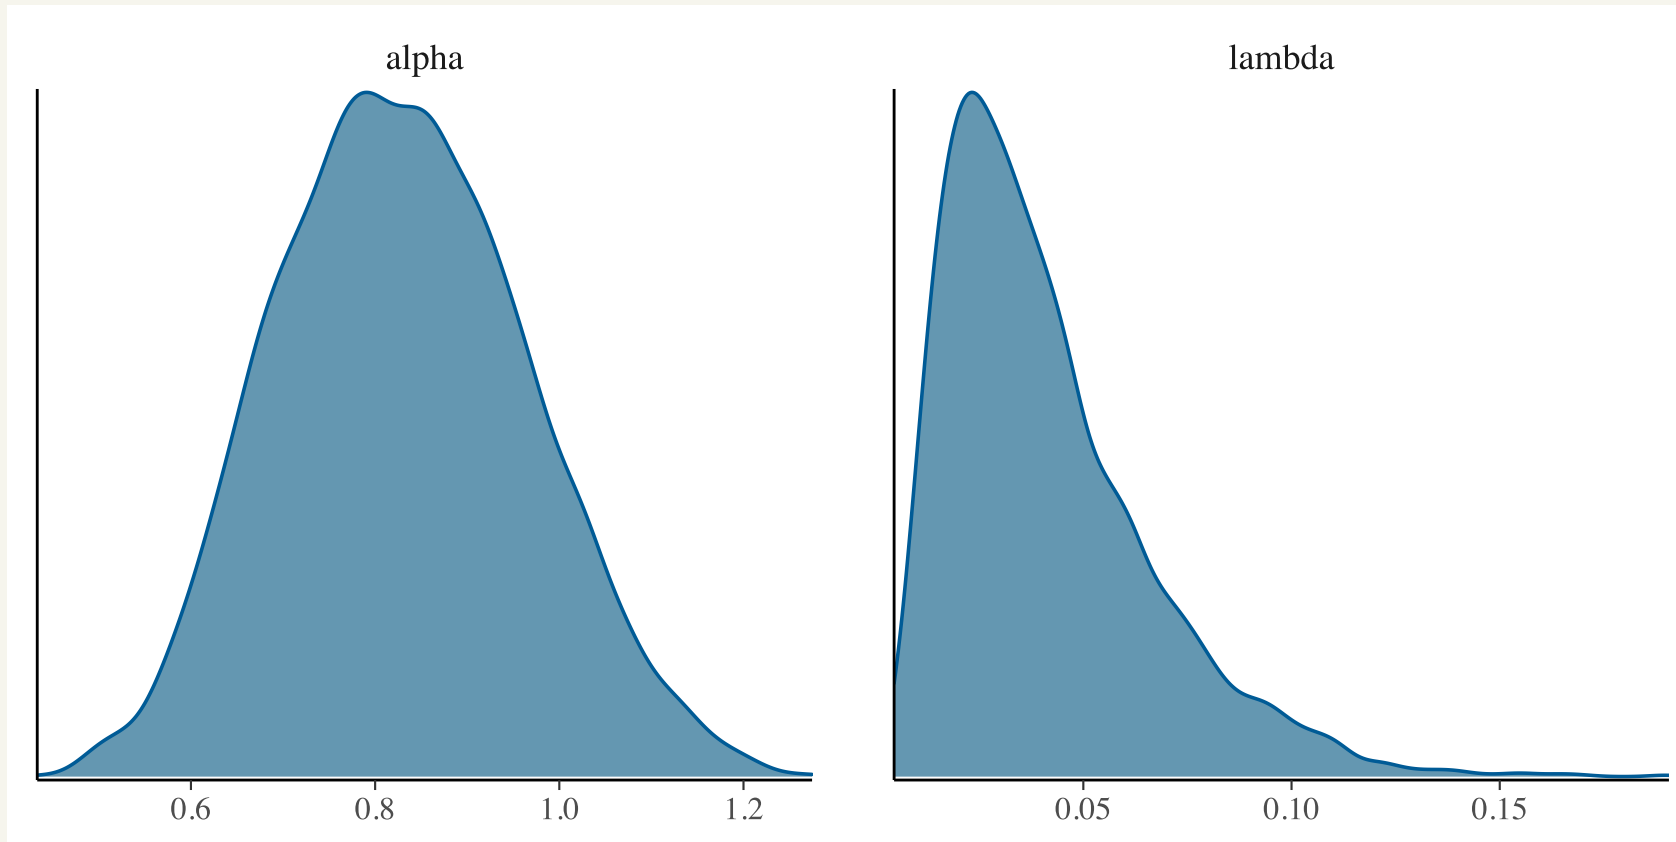
ACF plot

```
mcmc_acf(post_mcmc)
```



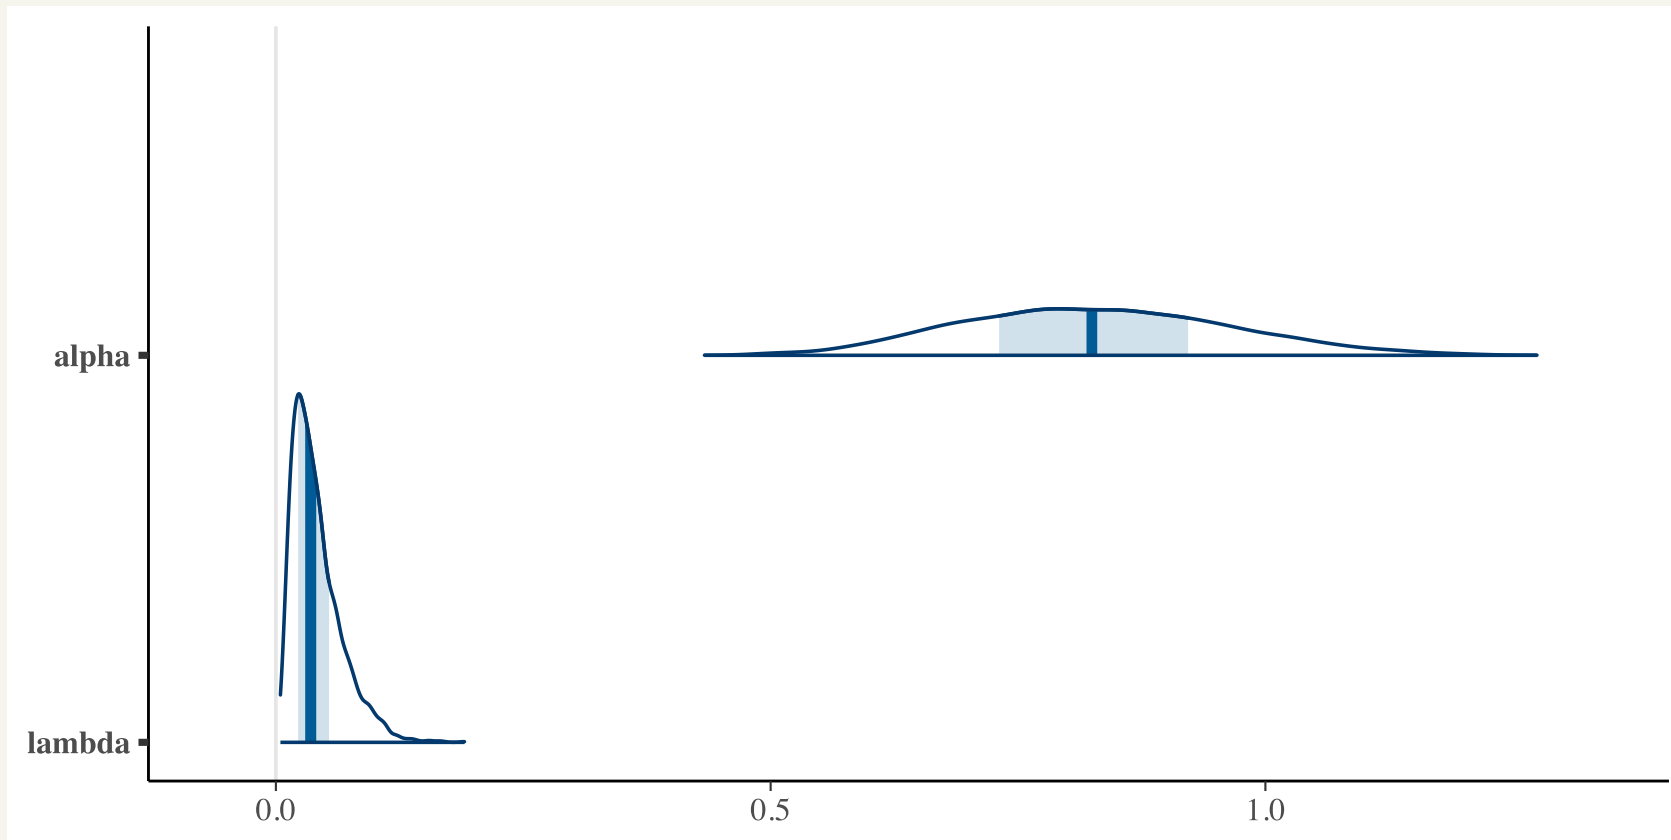
Posterior density plot

```
mcmc_dens(post_mcmc)
```



Plot equal-tail interval estimates

```
mcmc_areas(post_mcmc)
```



Your turn

The median survival time can be calculated as

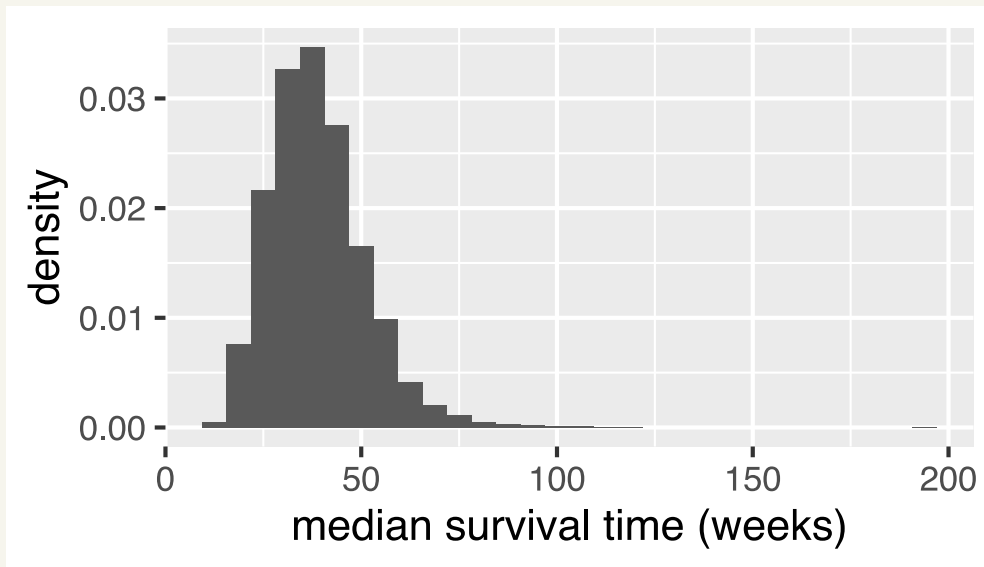
$$\left(\frac{\log(2)}{\lambda} \right)^{1/\alpha}$$

How could you calculate a 90% credible interval for the median survival time?

01:30

Posterior inference for `runjags` objects

```
median_survive <- (log(2) / posterior$mcmc[[1]][,"lambda"])^(1 / posterior$mcmc[[1]][,"alpha"])
```



90% credible interval for the median survival time

```
quantile(median_survive,  
         probs = c(0.05, 0.95))  
##          5%          95%  
## 21.82081 60.24946
```

Monitoring functions of parameters

You can create objects in JAGS using `<-`


```
model_string2 <- "  
model{  
  # Specify the likelihood  
  for(i in 1:n) {  
    y[i] ~ dweib(alpha, lambda)  
  }  
  
  # Specify the priors  
  alpha ~ dgamma(1, 1)  
  lambda ~ dgamma(1.53, 26.3)  
  
  # Calculate median  
  median <- (log(2) / lambda)^(1 / alpha)  
}  
"
```

Monitoring functions of parameters

You can monitor objects you create

```
posterior2 <- run.jags(  
  model = model_string2,  
  n.chains = 1,  
  data = survival_data,  
  monitor = "median",  
  adapt = 1000,  
  burnin = 5000,  
  sample = 5000  
)
```

Did the chain converge?



Your turn

Work through the example on the JAGS handout with a neighbor.

Check the website for a .Rmd template