

# More MCMC diagnostics

Stat 340: Bayesian Statistics

# Tuning MCMC

Three main decisions:

- Selecting the initial values for the parameters
- Determining if/when the chain(s) has converged
- Selecting the number of samples needed to approximate the posterior

# Initial values

- The algorithm will eventually converge no matter what initial values you select
- Choosing good starting values will speed up convergence
- It is important to try a few initial values to verify they all give the same result
- Usually 3-5 separate chains is sufficient

## **Common strategies:**

1. Select good initial values using method of moments or MLE
2. Purposely pick bad but different initial values for each chain to check convergence

# Convergence diagnostics


Don't assume that your Markov chain converged to the desired stationary distribution just because "it worked."

A **"good" Markov chain** should be:

- **Stationary:** draws should within the posterior (stable path around the center of gravity)
- **Well mixed:**
  - successive draws should not be highly correlated
  - each chain should target the same stationary distribution

# Your turn

**With a neighbor, decide whether each chain is well mixed**



# Tools already in our toolkit

Did my chains converge?

- Trace plots

Are my chains well mixed?

- ACF plots

# Convergence diagnostics

The **coda** R package has dozens of diagnostics

```
library(coda)
```

Did my chains converge?

- Geweke
- Gelman-Rubin

Did I run the sampler long enough after convergence?

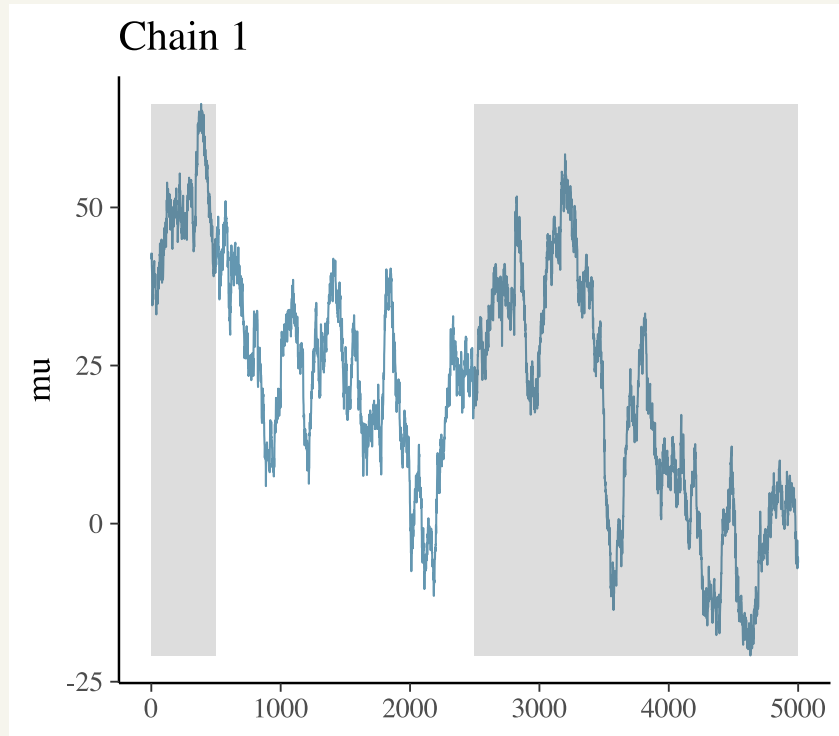
- Effective sample size
- Standard errors for the posterior mean estimate

# Geweke diagnostic

- **Idea:** a two-sample test comparing the mean of a chain between batches at the beginning versus the end
- By default, JAGS compares the first 10% with the last 50%
- Test statistic is a Z-score with the standard errors adjusted for autocorrelation (so we won't write down the formula)




# Geweke diagnostic



```
geweke.diag(bad_samples$mcmc[[1]])
```

```
##  
## Fraction in 1st window = 0.1  
## Fraction in 2nd window = 0.5  
##  
##      mu  
## 2.676
```

# Geweke diagnostic




```
geweke.diag(good_samples$mcmc[[2]])
```

```
##  
## Fraction in 1st window = 0.1  
## Fraction in 2nd window = 0.5  
##  
##      mu  
## 1.704
```

# Gelman-Rubin diagnostic

- Effective convergence of Markov chain simulation has been reached when inferences for quantities of interest do not depend on the starting point of the simulations.
- If we run multiple chains, we hope that all chains give same result
- Gelman and Rubin (1992) proposed (essentially) an ANOVA test of whether the chains have the same mean
- $R_j$  is scaled and approaches 1 from above
  - $R_j = 1 \Rightarrow$  perfect convergence
  - $R_j \geq 1.1 \Rightarrow$  red flag


# Gelman-Rubin diagnostic



```
gelman.diag(bad_samples$mcmc)
```

```
## Potential scale reduction factors
##
##      Point est. Upper C.I.
## mu           1.06      1.17
```

# Gelman-Rubin diagnostic



```
gelman.diag(good_samples$mcmc)
```

```
## Potential scale reduction factors  
##  
##      Point est. Upper C.I.  
## mu              1          1
```

# A note on autocorrelation

- Lower values are better, but if the chains are long enough even large values can be OK
- **Thinning** the Markov chain means keeping only every  $k$ th draw, where  $k$  is chosen so that the autocorrelation is small
- `thin` argument to `run.jags()` implements this
- This is always less efficient than using all samples, but can save memory


# Effective sample size

- Highly correlated samples have less information than independent samples
- $S = \#$  MCMC samples after burn in
- **Effective samples size (ESS)**

$$ESS_k = S / \left( 1 + 2 \sum_{k=1}^{\infty} \rho(k) \right)$$

- ESS = "equivalent number of independent observations"
- Should be at least a few thousand for all parameters

# Effective sample size




```
# ESS for single chain of mu1  
# n.iter = 5000  
effectiveSize(bad_samples$mcmc[[1]])
```

```
##          mu  
## 6.062577
```



# Effective sample size



```
# ESS for single chain of mu1  
# n.iter = 5000  
effectiveSize(good_samples$mcmc[[1]])
```

```
##          mu  
## 2119.896
```

# Standard errors of posterior mean estimates

- Assuming independence the standard error is

$$\text{Naive SE} = \frac{s}{\sqrt{S}}$$

where  $s$  = sample SD

- A more realistic standard error is

$$\text{Time-series SE} = \frac{s}{\sqrt{ESS}}$$

- If the SE is too large, rerun the MCMC algorithm for a larger number of samples

# Standard errors of posterior mean estimates

```
summary(bad_samples$mcmc)
##
## Iterations = 2001:7000
## Thinning interval = 1
## Number of chains = 3
## Sample size per chain = 5000
##
## 1. Empirical mean and standard deviation for each variable,
##    plus standard error of the mean:
##
##           Mean           SD      Naive SE Time-series SE
##          15.455          17.026          0.139          3.341
##
## 2. Quantiles for each variable:
##
##    2.5%    25%    50%    75%    97.5%
## -17.812    3.998  15.898  27.229  48.263
```

# Standard errors of posterior mean estimates

```
summary(good_samples$mcmc)
##
## Iterations = 2001:7000
## Thinning interval = 1
## Number of chains = 3
## Sample size per chain = 5000
##
## 1. Empirical mean and standard deviation for each variable,
##    plus standard error of the mean:
##
##           Mean           SD      Naive SE Time-series SE
##      -0.56146      1.26084      0.01029      0.01574
##
## 2. Quantiles for each variable:
##
##      2.5%      25%      50%      75%      97.5%
## -3.6330 -1.2123 -0.3600  0.3301  1.2965
```

# What to do if the chains haven't converged?

- Increase the number of iterations
- Tune the Metropolis candidate distribution
- Use better initial values
- Use a more advanced algorithm
- Simplify/reparameterize the model
- Use (more) informative priors