

# Bayesian regression

Stat 340: Bayesian Statistics

# Example

- Partial census data for the Dobe area !Kung San, a foraging population
- Compiled from Nancy Howell's interviews


##		mean	sd	5.5%	94.5%	histogra
##	height	154.59709	7.7423321	142.8750	167.00500	
##	weight	44.99049	6.4567081	35.1375	55.76588	
##	age	41.13849	15.9678551	20.0000	70.00000	
##	male	0.46875	0.4997328	0.0000	1.00000	



## Life Histories of the **DOBE !KUNG**

FOOD, FATNESS, AND WELL-BEING OVER THE LIFE-SPAN

NANCY HOWELL



How can we write a **Bayesian** model to relate weight and height of the Kalahari foragers?

# Observation-specific mean

We can adapt our normal model for the mean to use an observation-specific mean,  $\mu_i$ :

Sampling model:  $Y_i | \mu_i, \sigma \stackrel{\text{ind}}{\sim} \mathcal{N}(\mu_i, \sigma)$

Now we need to link  $\mu_i$  and  $x_i$

# A weakly informative prior

We may have limited prior information about the regression coefficients,  $\beta_0$  and  $\beta_1$ , and/or the standard deviation  $\sigma$

Assume independence:  $\pi(\beta_0, \beta_1, \sigma) = \pi(\beta_0, \beta_1)\pi(\sigma)$


1. Assume  $\beta_0$  and  $\beta_1$  are independent

- $\pi(\beta_0, \beta_1) = \pi(\beta_0)\pi(\beta_1)$
- Assign a weakly informative prior each coefficient:  $\beta_i \sim \mathcal{N}(m_i, s_i)$
- example:  $\mathcal{N}(0, 100)$

2. Assign a weakly informative prior to  $\sigma$

- example:  $1\sigma^2 \sim \text{Gamma}(1, 1)$

# Prior predictive as "sanity check"



## Simulate the prior predictive:

1. Draw parameters from their prior distributions
2. Draw data from the sampling model plugging in these simulated parameters

# Sampling from the prior predictive distribution

```
nsim <- 1e4                                     # no. of simulations

prior.sims <- data_frame(                       # simulate from ind. priors
  beta0 = rnorm(nsim, 178, 100),
  beta1 = rnorm(nsim, 0, 10),
  sigma = runif(nsim, 0, 50)
)
```

```
weight.value <- 45                             # condition on value of x
```

```
prior.pred <- prior.sims %>%
  mutate(
    mu = beta0 + beta1 * weight.value,          # calculate  $E(y|x)$ 
    height = rnorm(n(), mean = mu, sd = sigma) # draw sample from normal
  )
```

# Deriving the posterior

**Sampling model:**  $Y_i|x_i, \beta_0, \beta_1, \sigma \sim \mathcal{N}(\beta_0 + \beta_1 x_i, \sigma)$

**Joint likelihood:**

$$\begin{aligned} L(\beta_0, \beta_1, \sigma) &= \prod_{i=1}^n \left[ \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left\{ -\frac{1}{2\sigma^2} (y_i - \beta_0 - \beta_1 x_i)^2 \right\} \right] \\ &\propto \left( \frac{1}{\sigma^2} \right)^{n/2} \exp\left\{ -\frac{1/\sigma^2}{2} \sum_{i=1}^n (y_i - \beta_0 - \beta_1 x_i)^2 \right\} \end{aligned}$$

**Joint posterior:**  $\pi(\beta_0, \beta_1, \sigma|\text{data}) \propto \pi(\beta_0, \beta_1, \sigma)L(\beta_0, \beta_1, \sigma)$



# JAGS for Bayesian SLR

Write down the model string

```
slr_model <- "model {  
  ## sampling model  
  for (i in 1:N){  
    y[i] ~ dnorm(beta0 + beta1 * x[i], invsigma2)  
  }  
  
  ## priors  
  beta0 ~ dnorm(mu0, g0)  
  beta1 ~ dnorm(mu1, g1)  
  invsigma2 ~ dgamma(a, b)  
  sigma <- sqrt(pow(invsigma2, -1))  
}"
```

# JAGS for Bayesian SLR

Define the data and set prior parameters

```
the_data <- list(  
  y = adults$height,      # response variable  
  x = adults$weight,      # explanatory variable  
  N = nrow(adults),       # sample size  
  mu0 = 0,                # prior mean for beta0  
  g0 = 0.0001,            # prior precision for beta0  
  mu1 = 0,                # prior mean for beta1  
  g1 = 0.0001,            # prior precision for beta1  
  a = 1,                  # prior shape for 1/sigma2  
  b = 1,                  # prior scale for 1/sigma2  
)
```

# JAGS for Bayesian SLR

Generate samples from the posterior

```
posterior <- run.jags(  
  slr_model,  
  data = the_data,  
  n.chains = 1,  
  monitor = c("beta0", "beta1", "sigma"),  
  adapt = 1000,  
  burnin = 5000,  
  sample = 5000,  
  silent.jags = TRUE  
)
```


# Summary of the fitted model

```
summary(posterior)
```

```
##              Lower95      Median      Upper95      Mean      SD      Mode
## beta0 110.6005490 114.0250921 117.6466865 114.1046727 1.8346303 113.8677869
## beta1  0.8222891  0.9017146  0.9775988  0.9000988 0.0403727  0.9064987
## sigma  4.7315885  5.0787794  5.4687267  5.0857117 0.1917283  5.0777442
##              MCerr MC%ofSD SSeff      AC.10 psrf
## beta0 0.247707195      13.5    55 0.79203877    NA
## beta1 0.005459015      13.5    55 0.79115579    NA
## sigma 0.002711448       1.4 5000 0.01892964    NA
```


# MCMC diagnostics

```
mcmc_trace(posterior$mcmc)
```



# MCMC diagnostics

```
mcmc_acf(posterior$mcmc)
```




# Setting `thin = 50`

```
posterior <- run.jags(  
  slr_model,  
  data = the_data,  
  n.chains = 1,  
  monitor = c("beta0", "beta1", "sigma"),  
  adapt = 1000,  
  burnin = 5000,  
  sample = 5000,  
  thin = 50,  
  silent.jags = TRUE  
)
```

# MCMC diagnostics


```
mcmc_trace(posterior$mcmc)
```





# MCMC diagnostics

```
mcmc_acf(posterior$mcmc)
```




# Summary of the fitted model

```
summary(posterior)
```

```
##              Lower95      Median      Upper95      Mean      SD      Mod
## beta0 110.1966997 113.7337743 117.6820902 113.7703854 1.94096529 113.854709
## beta1  0.8216034  0.9082515  0.9880905  0.9073003 0.04272439  0.912202
## sigma  4.6997300  5.0820890  5.4454647  5.0856568 0.19255547  5.067288
##              MCerr MC%ofSD SSeff      AC.500 psrf
## beta0 0.0404676912      2.1  2300 0.008030463    NA
## beta1 0.0009197745      2.2  2158 0.006204312    NA
## sigma 0.0026375062      1.4  5330 -0.010200225    NA
```

# Plotting the fitted model


```
post_means <- apply(  
  posterior$mcmc[[1]], 2, mean  
)
```



# Sampling from the joint posterior


```
post_draws <- as.data.frame(  
  posterior$mcmc[[1]]  
)  
head(post_draws)
```

##		beta0	beta1	sigma
##	6001	116.1242	0.8590673	5.221635
##	6051	113.7029	0.9053650	4.954524
##	6101	111.2623	0.9591552	4.792757
##	6151	113.9549	0.9026097	4.987096
##	6201	117.3223	0.8261667	5.039425
##	6251	113.3996	0.9095483	5.626173



# Generating mean responses

```
mu_at_50 <- with(post_draws,  
                  beta0 + beta1 * 50)
```



```
quantile(mu_at_50, probs = c(0.05, 0.95))  
##           5%           95%  
## 158.5648 159.7029
```

